

Lecture 17: B-Trees (2-3, 2-3-4 Trees)

10/5/2020

BST Tree Height

BST Tree Height

- Trees range from best-case "bushy" to worst-case "spindly"
 - Height varies dramatically among the two
 - $\Theta(\log N)$ for bushy vs $\Theta(N)$ for spindly
- Performance of operations on spindly trees can be just as bad as a linked list!
- A worst case (spindly tree) has a height that grows exactly linearly - $\Theta(N)$
- A best case (bushy tree) has a tree height that grows exactly logarithmically - $\Theta(\log N)$

The Usefulness of Big O

- Big O is a useful idea:
 - Allows us to make simple blanket statements, e.g. can just say "binary search is $O(\log N)$ " instead of "binary search is $\Theta(\log N)$ in the worst case"
 - Sometimes don't know the exact runtime, so use O to give an upper bound
 - Example: Runtime for finding shortest route that goes to all world cities is $O(2^N)$. There might be a faster way, but nobody knows one yet
 - Easier to write proofs for Big O than Big Theta, e.g. finding runtime of mergesort, you can round up the number of items to the next power of 2. A little beyond the scope of the course.

Height, Depth, and Performance

Height and Depth

- Height and average depth are important properties of BSTs
 - The **"depth" of a node** is how far it is from the root
 - The root has depth 0
 - The **"height" of a tree** is the depth of its deepest leaf
 - The **"average depth"** of a tree is the average depth of a tree's nodes

Height, Depth, and Runtime

- Height and average depth determine runtimes for BST operations
 - The **"height"** of a tree determines the worst case runtime to find a node
 - The **"average depth"** determines the average case runtime to find a node

Important Question: What about real world BSTs?

- BSTs have:
 - Worst case $\Theta(N)$ height
 - Best case $\Theta(\log N)$ height

- One way to approximate real world BSTs is to consider randomized BSTs
- **Nice Property.** Random trees have $\Theta(\log N)$ average depth and height
 - In other words: Random trees are bushy, not spindly

Randomized Trees: Mathematical Analysis

- **Average Depth.** If N distinct keys are inserted into a BST, the expected average depth is $\sim 2 \ln N$
 - Thus, average runtime for contains operation is $\Theta(\log N)$ on a tree built with random inserts
- **Tree Height.** If N distinct keys are inserted in random order, expected tree height is $\sim 4.311 \ln N$
 - Thus, worst case runtime for contains operation is $\Theta(\log N)$ on a tree built with random inserts
- BSTs have:
 - Worst case $\Theta(N)$ height
 - Best case $\Theta(\log N)$ height
 - $\Theta(\log N)$ height if constructed via random inserts
- In real world applications we expect both insertion and deletion
 - Can show that random trees including deletion are still $\Theta(\log N)$ height


Good News and Bad News

- Good news: BSTs have great performance if we insert items randomly
 - Performance is $\Theta(\log N)$ per operation
- Bad news: We can't always insert our items in a random order
 - Data comes in over time, don't have all at once

B-trees / 2-3 trees / 2-3-4 trees

Avoiding Imbalance through Overstuffing

- The problem is adding new leaves at the bottom
- Crazy idea: never add new leaves at the bottom
 - Tree can never get imbalanced
- Avoid new leaves by "overstuffing" the leaf nodes
 - "Overstuffed tree" always has balanced height, because leaf depths never change
- Overstuffed trees are a logically consistent but very weird data structure




- contains(18):
 - $18 > 13$? Yes, go right

- $18 > 15$? Yes, go right
- $16 = 18$? No
- $17 = 18$? No
- $18 = 18$? Yes! Found it
- Problem with this idea? Degenerates into linked list


Revising Our Overstuffed Tree Approach: Moving Items Up

- Height is balanced, but we have a new problem
 - Leaf nodes can get too juicy
- Solution?
 - Set a limit L on the number of items, say $L=3$
 - If any node has more than L items, give an item to parent
 - Which one? Let's say (arbitrarily) the left-middle
- What's the problem now?
 - 16 is to the right of 17



Revising Overstuffed Tree Approach: Node Splitting

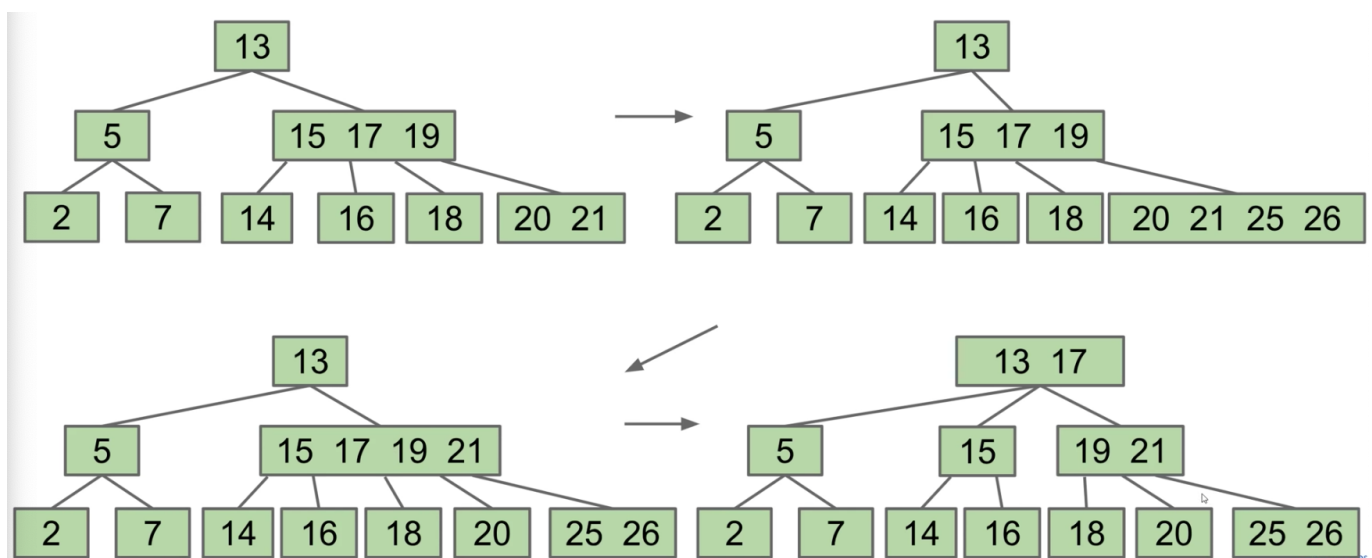
- Solution?
 - Set a limit L on the number of items, say $L=3$
 - If any node has more than L items, give an item to parent
 - Pulling item out of full node splits it into left and right
 - Parent node now has three children!



- This is a logically consistent and not so weird data structure
 - Contains(18):
 - $18 > 13$, so go right
 - $18 > 15$, so compare vs. 17
 - $18 > 17$, so go right
- Examining a node costs us $O(L)$ compares, but that's OK since L is constant
- What if a non-leaf node gets too full? Can we split that?

add: Chain Reaction


- Suppose we add 25, 26:



What Happens if the root is too full?



Challenge: Draw the tree after the root is split.



Perfect Balance

- Observation: Splitting-trees have perfect balance
 - If we split the root, every node gets pushed down by exactly one level
 - If we split a leaf or internal node, the height doesn't change
- All operations have guaranteed $O(\log N)$ time

The Real Name for Splitting Trees is "B Trees"

- B-trees of order $L=3$ (like we used today) are also called a 2-3-4 tree or a 2-4 tree
 - "2-3-4" refers to the number of children that a node can have
- B-trees of order $L=2$ are also called a 2-3 tree
- B-Trees are most popular in two specific contexts:
 - Small L ($L=2$ or $L=3$)
 - Used as a conceptually simple balanced search tree
 - L is very large (say thousands)
 - Used in practice for databases and file systems

B-tree Bushiness Invariants

Exercise

- No matter the insertion order you choose, resulting B-Tree is always bushy!
 - May vary in height a little bit, but overall guaranteed to be bushy

B-Tree Invariants

- Because of the way B-Trees are constructed, we get two nice invariants
 - All leaves must be the same distance from the source
 - A non-leaf node with k items must have exactly $k+1$ children
- These invariants guarantee that our tree will be bushy

B-Tree Runtime Analysis

Height of a B-Tree with Limit L

- L: Max number of items per node
- Height: Between $\sim \log_{L+1}(N)$ and $\sim \log_2(N)$
 - Largest possible height is all non-leaf nodes have 1 item
 - Smallest possible height is all nodes have L items
 - Overall height is therefore $\Theta(\log N)$

Runtime for **contains**

- Runtime of contains:
 - Worst case number of nodes to inspect: $H + 1$
 - Worst case number of items to inspect per node: L
 - Overall runtime: $O(HL)$
- Since $H = \Theta(\log N)$, overall runtime is $O(L \log N)$
 - Since L is a constant, runtime is therefore $O(\log N)$
- Bottom line: contains and add are both $O(\log N)$

Summary

Summary

- BSTs have best case height $\Theta(\log N)$ and worst case height $\Theta(N)$
 - Big O is not the same thing as worst case
- B-Trees are a modification of the binary search tree that avoids $\Theta(N)$ worst case
 - Nodes may contain between 1 and L items
 - **contains** works almost exactly like a normal BST
 - **add** works by adding items to existing leaf nodes
 - If nodes are too full, they split
 - Resulting tree has perfect balance. Runtime for operations is $O(\log N)$
 - Have not discussed deletion
 - Have not discussed how splitting works if $L > 3$
 - B-trees are more complex, but they can efficiently handle ANY insertion order