

Lecture 19: Multidimensional Data

10/9/2020

Range-Finding and Nearest

Search Trees

- We've seen three different implementations of a Map
 - BST
 - 2-3 / B-Tree
 - Red Black Tree
- "search tree" data structures support very fast insert, remove, and delete operations for arbitrary amounts of data
 - Requires that data can be compared to each other with some total order
 - We used the "Comparable" interface as our comparison engine

Expanding the Power of our Set

- There are other operations we might want to include:
 - `select(int i)`: Returns the *i*th smallest element in the set
 - `rank(T x)`: Returns the "rank" of *x*
 - `subSet(T from, T to)`: returns all items between *from* and *to*
 - `nearest(T x)`: Returns the value closest to *x*

Implementing Fancier Set Operations with a BST

- It turns out that a BST can efficiently support the `select`, `rank`, `subSet`, and `nearest`
- How would you find `nearest(N)`?
 - Just search for *N* and record closest item seen
 - Exploits the BST structure

Sets and Maps on 2D Data

- So far we've only discussed "one dimensional data". That is, all data could be compared under some total order
- But not all data can be compared along a single dimension
 - We'll see that search trees require some design tweaks to function efficiently on multi-dimensional data

Multi-Dimensional Data

Motivation: 2D Range Finding and Nearest Neighbors

- Suppose we want to perform operations on a set of Body objects in 2D space
 - 2D range searching: How many objects are in a highlighted rectangle
 - Nearest: What is the closest object?

- Ideally, we'd like to store our data in a format that allows more efficient approaches than just iterating over all objects
- It's difficult to build a BST of 2 dimensional data
 - Difficult to compare objects, lose some information about ordering

QuadTrees

The QuadTree

- A QuadTree is the simplest solution conceptually
 - Every Node **four** children
 - Top left (northwest)
 - Top right (northeast)
 - Bottom left (southwest)
 - Bottom right (southeast)
- Just like a BST, insertion order determines the topology of a QuadTree

QuadTrees

- Quadtrees are a form of "spatial partitioning"
 - Quadtrees: Each node "owns" 4 subspaces
 - Space is more finely divided in regions where there are more points
 - Results in better runtime in many circumstances

QuadTree Range Search

- Quadtrees allow us to prune when performing a rectangle search
 - Simple idea: Prune subspaces that don't intersect the query rectangle

Higher Dimensional Data

3D Data

- Suppose we want to store objects in 3D space
 - Quadtrees have only four directions, but in 3D, there are 8
- One approach: Use an Oct-tree or Octree
 - Very widely used in practice

Even Higher Dimensional Space

- You may want to organize data on a larger number of dimensions
- In these cases, one somewhat common solution is a k-d tree
 - Fascinating data structure that handles arbitrary numbers of dimensions
 - k-d means "k dimensional"
 - For the sake of simplicity, we'll use 2D data, but the idea generalizes naturally

K-d Trees

- k-d tree example for 2-d
 - Basic idea, root node partitions entire space into left and right (by x)

- All depth 1 nodes partition subspace into up and down (by y)
- All depth 2 nodes partition subspace into left and right (by x)
- And continue alternating down the depth of the tree
- To break ties, we'll say items that are equal in one dimension go off to the right (or up) child of each node
- Each point owns 2 subspaces
 - Similar to a quadtree

K-d Trees and Nearest Neighbor

- You can simplify code by only measuring the length of vertical/horizontal lines instead of diagonal hypotenuses.
 - Optimization: Do not explore subspaces that can't possibly have a better answer than your current best

Nearest Pseudocode

- nearest(Node n, Point goal, Node best):
 - If n is null, return best
 - If $n.\text{distance}(\text{goal}) < \text{best}.\text{distance}(\text{goal})$, best = n
 - If goal < n (according to n's comparator):
 - goodSide = n."left"Child
 - badSide = n."right"Child
 - else:
 - goodSide = n."right"Child
 - badSide = n."left"Child
 - best = nearest(goodSide, goal, best)
 - If bad side could still have something useful
 - best = nearest(badSide, goal, best)
 - return best

Uniform Partitioning

Uniform Partitioning

- Not based on a tree at all
- Simplest idea: Partition space into uniform rectangular buckets (also called "bins")
- Algorithm is still $\Theta(N)$, but it's faster than iterating over all the points

Uniform vs. Hierarchical Partitioning

- All of our approaches today boil down to spatial partitioning
 - Uniform partitioning (perfect grid of rectangles)
 - Quadtrees and KdTrees: Hierarchical partitioning
 - Each node "owns" 4 and 2 subspaces, respectively
 - Space is more finely divided into subspaces when there are more points

Uniform Partitioning vs. Quadtrees and Kd-Trees

- Uniform partitioning is easier to implement than a QuadTree or Kd-Tree
 - May be good enough for many applications

Summary and Applications

Multi-Dimensional Data Summary

- Multidimensional data has interesting operations:
 - **Range Finding**
 - **Nearest**
- The most common approach is **spatial partitioning**:
 - **Uniform Partitioning**: Analogous to hashing
 - **QuadTree**: Generalized 2D BST where each node "owns" 4 subspaces
 - **K-d Tree**: Generalized k-d BST where each node "owns" 2 subspaces
 - Dimension of ownership cycles with each level of depth in tree
- Spatial partitioning allows for **pruning** of the search space