

Lecture 25: Shortest Paths

10/23/2020

Graph Problems

- Which is better?
 - DFS or BFS

BFS vs. DFS for Path Finding

- Possible considerations:
 - **Correctness.** Do both work for all graphs?
 - Yes
 - **Output Quality.** Does one give better results?
 - BFS is a 2-for-1 deal, not only do you get paths, but your paths are also guaranteed to be shortest
 - **Time Efficiency.** Is one more efficient than the other?
 - Should be very similar. Both consider all edges twice
 - **Space Efficiency.** Is one more efficient than the other?
 - DFS is worse for spindly graphs
 - Call stack gets very deep
 - Computer needs $\Theta(V)$ memory to remember recursive calls
 - BFS is worse for absurdly "bushy" graphs
 - Queue gets very large. In worst case, queue will require $\Theta(V)$ memory
 - Note: In our implementations, we have to spend $\Theta(V)$ memory anyway to track `distTo` and `edgeTo` arrays

Breadth FirstSearch for Google Maps

- BFS would not be a good choice for a google maps style navigation application
 - We need an algorithm that takes into account edge distances, also known as "edge weights"

Dijkstra's Algorithm

Single Source Single Target Shortest Paths

- Observation: Solution will always be a path with no cycles (assuming non-negative weights)

Problem: Single Source Shortest Paths

- Goal: Find the shortest paths from **source** vertex s to every other vertex
- Observation: Solution will always be a **tree**
 - Can think of as the union of the shortest paths to all vertices

Edge Count

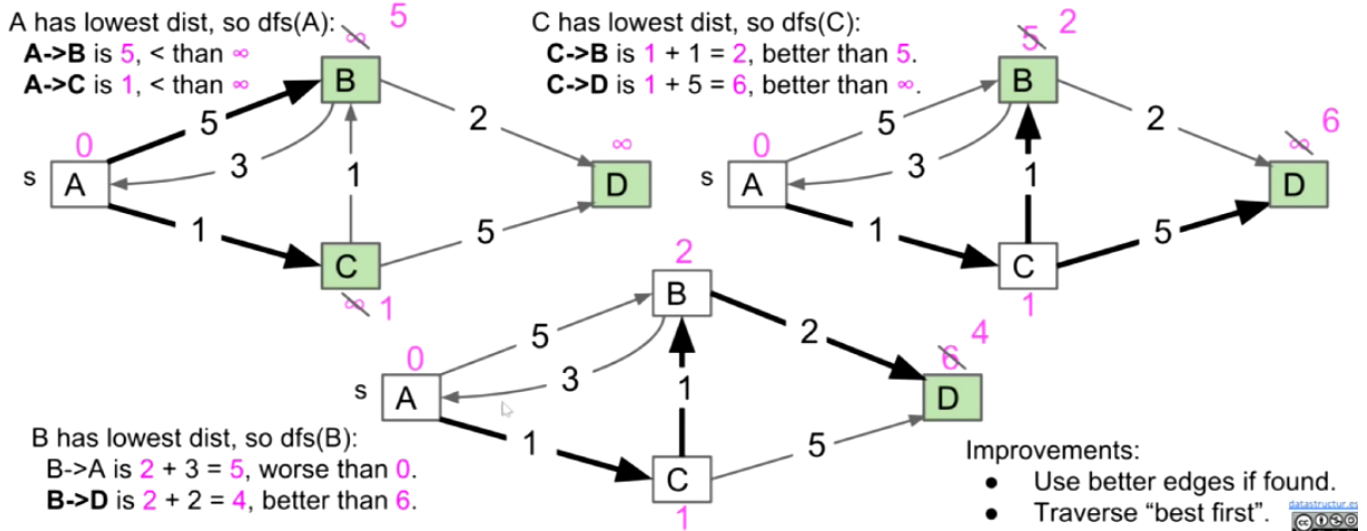
- If G is a connected edge-weighted graph with V vertices and E edges, how many edges are in the **Shortest Paths Tree** (SPT) of G ? [assume every vertex is reachable]
 - Since the solution is a tree, there are $V-1$ edges

Creating an Algorithm


- Start with a bad algorithm
 - Algorithm begins with all vertices unmarked and all distance infinite. No edges in the shortest paths tree (SPT)
- Bad algorithm #1: Perform a depth first search. When you visit v :
 - For each edge from v to w , if w is not already part of SPT, add the edge
 - Note: This WILL NOT WORK
- Bad algorithm #2: Perform a depth first search. When you visit v :
 - For each edge from v to w , add edge to the SPT **only if that edge yields better distance** (we'll call this process "edge **relaxation**")
 - Improvements:
 - Use better edges if found

Dijkstra's Algorithm

- Perform a **best first search** (closest first). When you visit v :
 - For each v to w , **relax that edge**
 - Improvements:
 - Use better edges if found
 - Traverse "best first"
 - Insert all vertices into fringe PQ (e.g. use a heap), storing vertices in order of distance from source
 - Repeat: Remove (closest) vertex v from PQ, and relax all edges pointing from v
 - Note: If non-negative weights, **impossible for any inactive vertex (i.e. already visited and not on the fringe) to be improved**
 - Would result in a cycle if it does



#	distTo	edgeTo
0	0	-
1	2	0
2	1	0
3	11	6
4	5	1
5	9	4
6	10	4



Fringe: [(3: 11)]

Dijkstra's Correctness and Runtime

Dijkstra's Algorithm Pseudocode

- Dijkstra's:
 - PQ.add(source, 0)
 - For other vertices v , PQ.add(v , infinity)
 - While PQ is not empty:
 - p = PQ.removeSmallest()
 - Relax all edges from p
- Relaxing** and edge $p \rightarrow q$ with weight w :
 - If $\text{distTo}[p] + w < \text{distTo}[q]$:
 - $\text{distTo}[q] = \text{distTo}[p] + w$
 - $\text{edgeTo}[q] = p$
 - PQ.changePriority(q , $\text{distTo}[q]$)
- Key invariants:
 - $\text{edgeTo}[v]$ is the best known predecessor of v
 - $\text{distTo}[v]$ is the best known total distance from source to v
 - PQ contains all unvisited vertices in order of distTo
- Important properties:

- Always visits vertices in order of total distance from source
- Relaxation always fails on edges to visited (white) vertices

Guaranteed Optimality

- Dijkstra's Algorithm
 - Visit vertices in **order of best-known distance** from source. On visit, **relax** every edge from the visited source
- Guaranteed to return a correct result if all edges are non-negative
 - Proof relies on the property that relaxation always fails on edges to visited vertices
- Proof sketch: Assume all edges have non-negative weights
 - At start, $\text{distTo}[\text{source}] = 0$, which is optimal
 - After relaxing all edges from source, let vertex v_1 be the vertex with minimum weight, i.e. that is closest to the source. Claim: $\text{distTo}[v_1]$ is optimal, and thus future relaxations will fail. Why?
 - $\text{distTo}[p] \geq \text{distTo}[v_1]$ for all p , therefore
 - $\text{distTo}[p] + w \geq \text{distTo}[v_1]$
 - Can use induction to prove that this holds for all vertices after dequeuing

Negative Edges

- Dijkstra's Algorithm
 - Visit vertices in **order of best-known distance** from source. On visit, **relax** every edge from the visited vertex
- Dijkstra's can fail if graph has negative weight edges
 - Relaxation of already visited edges can succeed

Dijkstra's Algorithm Runtime

- Priority Queue operation count, assuming binary heap based PQ:
 - add: V , each costing $O(\log V)$ time
 - removeSmallest: V , each costing $O(\log V)$ time
 - changePriority: E , each costing $O(\log V)$ time
- Overall runtime: $O(V \cdot \log(V) + V \cdot \log(V) + E \cdot \log(V))$
 - Assuming $E > V$, this is just $O(E \log V)$ for a connected graph

A*

Single Target Dijkstra's

- Is this a good algorithm for a navigation application
 - Will it find the shortest path?
 - Yes!
 - Will it be efficient
 - No. It will look for shortest path to other places

The Problem with Dijkstra's

- We have only a **single target** in mind, so we need a different algorithm. How can we do better?

How can we do better?

- Explore one direction first?

Introducing A*

- Simple idea:
 - Visit vertices in order of $d(\text{Denver}, v) + h(v, \text{goal})$, where $h(v, \text{goal})$ is an estimate of the distance from v to our goal NYC
 - In other words, look at some location if:
 - We already know the fastest way to reach v
 - AND we suspect that v is also the fastest way to NYC taking into account the time to get to v
- Observations:
 - Not every vertex gets visited
 - Result is not a shortest paths tree for a vertex, but that's OK since we only care about a path to a single vertex

A* Heuristic Example

- How do we get our estimate?
 - Estimate is an arbitrary **heuristic** $h(v, \text{goal})$
 - heuristic: "using experience to learn and improve"
 - Doesn't have to be perfect

A* Heuristics (Not covered in this class)

Heuristics and Correctness

- For our version of A* to give the correct answer, our A heuristic must be:
 - **Admissible:** $h(v, \text{NYC}) \leq \text{true distance from } v \text{ to NYC}$
 - **Consistent:** For each neighbor w of v :
 - $h(v, \text{NYC}) \leq \text{dist}(v, w) + h(w, \text{NYC})$
 - Where $\text{dist}(v, w)$ is the weight of the edge from v to w

Consistency and Admissibility (Beyond scope)

- All consistent heuristics are admissible
 - "Admissible" means that the heuristic never overestimates

Summary: Shortest Paths Problems

- Single source, multiple targets:
 - Can represent shortest path from start to every vertex as a shortest paths tree with $V-1$ edges
 - Can find the SPT using Dijkstra's algorithm
- Single source, single target:
 - Dijkstra's is inefficient (searches useless parts of the graph)
 - Can represent shortest path as path (with up to $V-1$ vertices, but probably far fewer)
 - A* is potentially much faster than Dijkstra's

- Consistent heuristic guarantees correct solution